

Atlas-PS 4

David Atlas

9/21/2018

Problem 1

We implement a Monte Carlo Integration function and apply it to $f(x) = 4\sqrt{1-x^2}, x \in [0, 1]$.

```
set.seed(73)

range0 <- function(avector){
  return(max(avector) - min(avector))
}

monte_carlo_integration <- function(f, a, b, n, plot=FALSE){

  # Get the range of the function on the space
  f_range <- range(f(seq(a, b, .01)))

  total_area <- range0(f_range) * (b - a)


  x_points <- runif(n, a, b)
  y_points <- runif(n, f_range[1], f_range[2])

  proportion_under <- mean(f(x_points) >= y_points)

  # Plot the scatter plot of points
  if(plot){
    plot(seq(a, b, .01), f(seq(a, b, .01)), 'l',
         lwd=5, col='blue', xlab='X', ylab=TeX("$f(x)$"))
    points(x_points, y_points,
           col= as.factor(ifelse(f(x_points) <= y_points, 0, 1)),
           lwd=.25)
  }

  return(total_area * proportion_under)
}

f <- function(x) 4 * sqrt(1 - x ^ 2)
monte_carlo_integration(f, 0, 1, 50000, plot=TRUE)
```



```
## [1] 3.14528
```

The solution is 3.14. This is as expected, as the integral defines the 4 times the area of a quarter unit circle, which has an area equal to $\frac{\pi}{4}$, and thus, the integral approximates the value of π .

Problem 2

a)

The PDF is

$$f(x) = \begin{cases} 4x & 0 < x \leq \frac{1}{2} \\ 4 - 4x & \frac{1}{2} \leq x \leq 1. \end{cases}$$

The CDF can be found by integrating each of the piecewise components. We subtract 1 from the second term to enforce continuity of the CDF, and to ensure that $F(\infty) = 1$

$$F(x) = \begin{cases} 2x^2 & 0 < x \leq \frac{1}{2} \\ 4x - 2x^2 - 1 & \frac{1}{2} < x \leq 1. \end{cases}$$


We plot the PDF and the CDF below.

```
pdf <- function(x) ifelse(x <= .5, 4 * x, 4 - 4*x)
cdf <- function(x) ifelse(x <= .5, 2 * x ^2, 4*x - 2*x^2 -1)

x <- seq(0, 1, .01)


plot(x, pdf(x), 'l', xlab='X', ylab=TeX('$f(x)$'), main='Plot of the PDF of X')
```

Plot of the PDF of X



```
plot(x, cdf(x), 'l', xlab='X', ylab=TeX('$F(x)$'), main='Plot of the CDF of X')
```

Plot of the CDF of X



b)

We find the inverse of the CDF. The inverse of the first piecewise component is derived first:

$$\begin{aligned}x &= 2F^{-1}(x)^2 \\ \implies F^{-1}(x) &= \sqrt{\frac{x}{2}}.\end{aligned}$$

The inverse of the second piecewise component is derived second. Here, we complete the square to isolate $F^{-1}(x)$.

$$\begin{aligned}x &= 4F^{-1}(x) - 2F^{-1}(x) - 1 \\ \implies -\frac{x}{2} &= F^{-1}(x) - 2F^{-1}(x) + \frac{1}{2} \\ \implies -\frac{x}{2} + \frac{1}{2} &= F^{-1}(x) - 2F^{-1}(x) + 1 = (F^{-1}(x) - 1)^2 \\ \implies F^{-1}(x) &= 1 - \sqrt{-\frac{x}{2} + \frac{1}{2}}.\end{aligned}$$

Additionally, we must ensure that the domain of the inverse CDF is appropriate; however, for this example, the domains actually remain the same. Additionally, the square roots can be either positive or negative, but due to the domain, $[0, 1]$, we can disregard the terms that are not relevant.


Therefore, we can write the inverse CDF as

$$F^{-1}(x) = \begin{cases} \sqrt{\frac{x}{2}} & 0 < x \leq \frac{1}{2} \\ 1 - \sqrt{-\frac{x}{2} + \frac{1}{2}} & \frac{1}{2} < x \leq 1. \end{cases}$$

c)

We generate 1000 random points using part b and plot a histogram.

```
set.seed(73)
U <- runif(1000, 0, 1)
G <- function(u) ifelse(u <= .5, sqrt(u / 2), 1 - sqrt(-.5 * u + .5) )
hist(G(U), breaks = seq(0, 1, .05))
```



```
inverse_cdf_sample <- G(U) # save this for later
```

The distribution converges on the original density function.

d)

We use rejection sampling to generate a random sample from the distribution. We can use a uniform distribution as the instrumental distribution, ($g(x) \sim U(0, 2)$) with an envelope $e(x) = \frac{g(x)}{1}$; $\alpha = 1$.


```
set.seed(73)

f <- function(x) ifelse(x <= .5, 4 * x, 4 - 4*x)

# Generate large samples so there will be 1000 non-rejected points
Y <- runif(100000, 0, 2)
U <- runif(100000, 0, 1)

# Get the first 1000 points that were not rejected
x <- U[Y < f(U)] [0:1000]
hist(x, breaks = seq(0, 1, .05), main="Rejection Sampling - Distribution of X", xlab='X')
```

Rejection Sampling – Distribution of X



Here, we see that the distribution is very similar (and appears to be converging) on the original density function.

e)

To use the Monte Carlo method to approximate $E[X^2]$, we take our sample from above and take the mean of the squared values. This is because we are trying to find $\int_0^1 g(x)f(x)dx$ where $f(x)$ is the piecewise density found above, and $g(x) = x^2$.

```
# Use the sample from above. Square the values and take the mean.
round(mean(inverse_cdf_sample^2), 3)
```

```
## [1] 0.286
```

Our simulation method yields .286. We need to find the actual value. We can write

$$\begin{aligned}
 E[X^2] &= \int_0^1 x^2 f(x) dx \\
 &= \int_0^{\frac{1}{2}} 4x^3 dx + \int_{\frac{1}{2}}^1 (4x^2 - 4x^3) dx \\
 &= \frac{1}{16} + \frac{1}{3} - \left(\frac{1}{6} - \frac{1}{16}\right) \\
 &= \frac{1}{16} + \frac{22}{96} = .29166.
 \end{aligned}$$


Therefore, our approximation is reasonably good. If we expanded the sample size, we could increase the accuracy of the estimate.

Problem 3

a)

We plot the two distributions given:

```
plot(seq(-10, 10, .01), dnorm(seq(-10, 10, .01), 0, 1), 'l',  
     ylim=c(0, .5),  
     main = "N(0, 1) and N(1, 2) with alpha=.4",  
     xlab = 'X', ylab='f(x)')  
lines(seq(-10, 10, .01), dnorm(seq(-10, 10, .01), 1, 2) / .4, 'l')
```




We can see that using $\alpha = .4$ satisfies our requirements for rejection sampling.

Next we generate a sample of 10000 points using rejection sampling.

```
n <- 0  
total <- 0  
X <- rep(0, 10000)  
while (n < 10000){  
  total <- total + 1  
  Y <- rnorm(1, 1, 2)  
  U <- runif(1, 0, 1)  
  if (U <= dnorm(Y, 0, 1) / (dnorm(Y, 1, 2) / .4)){  
    n <- n + 1  
    X[n] = Y  
  }  
}  
hist(X, breaks = seq(-5,5,.25))
```

Histogram of X



```
print(paste0("It took ", total,
             " random draws to accumulate 10000 points, leading to efficiency ",
             round(n / total, 3)))
```

```
## [1] "It took 24684 random draws to accumulate 10000 points, leading to efficiency 0.405"
```

```
print(paste0("The mean of the sample is: ", mean(X)))
```

```
## [1] "The mean of the sample is: -0.00670776807783822"
```

```
print(paste0("The standard deviation of the sample is: ", sd(X)))
```

```
## [1] "The standard deviation of the sample is: 1.01795364958788"
```

We see that the mean is almost exactly 0 and the standard deviation is almost exactly 1.

b)

Our choice for α was .4, which we derived by plotting the two functions and ensuring that the instrumental function was greater than the target function when divided by .4. Our envelope is therefore $e(x) = g(x)/.4$ where $g(x)$ is the density function for a normally distributed variable with mean 1 and standard deviation 2.

c)

Generally, squeezed rejection sampling can help reduce the number of times that $f(x)$ needs to be computed. If it is expensive to compute, this can save time, as the squeezing function can be easy to compute. The squeezing function simply acts as a filter to accept points that are easy to determine as being less than $f(x)$ without computing $f(x)$.

There isn't really any reason to use squeezed rejection sampling in this problem, as $f(x)$ is not computationally expensive to compute. Using the built in R function, it's a constant time evaluation.

However, if we were to want to use a squeezing function for this problem, we could use a piecewise function similar to the one in problem 2.


$$s(x) = \begin{cases} .38 + .2x & x < 0 \\ .38 - .2x & x \geq 0 \end{cases}$$

We plot the densities again with the squeezing function.

```
plot(seq(-10, 10, .01), dnorm(seq(-10, 10, .01), 0, 1), 'l',
     ylim=c(0, .5),
     main = "N(0, 1) and N(1, 2) with alpha=.4 and a squeezing function.",
     xlab = 'X', ylab='f(x)')
lines(seq(-10, 10, .01), dnorm(seq(-10, 10, .01), 1, 2) / .4, 'l')

s <- function(x) ifelse(x < 0, .38 + .2 * x, .38 - .2 * x)
lines(seq(-10, 10, .01), s(seq(-10, 10, .01)))
```

N(0, 1) and N(1, 2) with alpha=.4 and a squeezing function.



This could help, as we can evaluate the squeezing function easily, and use it to filter observations before evaluation $f(x)$.

Just for fun, I ran a quick time trial across a handful of values of n (the number of time evaluating each of these functions).

```
n <- c(10, 100, 1000, 10000, 100000)
norm_times <- rep(0, length(n))
s_times <- rep(0, length(n))
for (i in 1:length(n)){
  x <- runif(n[i], -1, 1)
  norm_begin <- Sys.time()
  norm <- dnorm(x, mean=0, sd=1)
  norm_end <- Sys.time()
```

```

s_begin <- Sys.time()
s_ <- s(x)
s_end <- Sys.time()

norm_times[i] <- norm_end - norm_begin
s_times[i] <- s_end - s_begin
}
options(scipen = 1000000)
knitr::kable(data.frame(n=n, f_times = norm_times, s_times = s_times),
              col.names = c("n", "Time to Evaluate f(x)", "Time to Evaluate s(x)"))

```

n	Time to Evaluate f(x)	Time to Evaluate s(x)
10	0.0000079	0.0021620
100	0.0000081	0.0000219
1000	0.0000310	0.0001349
10000	0.0002561	0.0012081
100000	0.0023060	0.0217569

In practice, it ends up being faster to compute $f(x)$ for large samples, so we wouldn't use the squeezing function here. Not too surprising, as $f(x)$ here is not computationally expensive.